

Dashboard | Amazon SageMaker

us-west-2.console.aws.amazon.com/sagemaker/home?region=us-west-2#/dashboard

Services

Search

[Alt+S]

Oregon

voclabs/user3044095=cf0b5626-590a-11ee-a8a3-13a98f4cdb72 @ 08...

Amazon SageMaker

Getting started

Studio

Studio Lab

Canvas

RStudio

TensorBoard

Profiler

Admin configurations

Domains

Role manager

Images

Lifecycle configurations

SageMaker dashboard

Search

JumpStart

Foundation models

Computer vision models

Natural language processing models

Governance

HyperPod Clusters

Ground Truth

Amazon SageMaker > Dashboard

Amazon SageMaker Studio

Explore Amazon SageMaker Studio, a machine learning Integrated Development Environment (IDE) to build, train and debug models, track experiments, deploy models in production, and monitor drifts.

Learn more

Dashboard

Open SageMaker Domain

Recent activity

Recent activity within the

Last 7 days

Ground Truth	Notebook	Training	Inference	Processing	Canvas
Labeling jobs	Notebook instances	Training jobs	Models	Processing jobs	Endpoints
No recent activity.	No recent activity.	2 Completed	No recent activity.	No recent activity.	No recent activity.
		2 Created	Endpoints		
		Hyperparameter tuning jobs	No recent activity.		
		No recent activity.	Batch transform jobs		
			No recent activity.		

Learning Content

Amazon SageMaker 10-Minute Studio Tutorial

Step-by-step guide to getting started with Studio faster. Learn more

Amazon SageMaker How-to Blog

AWS machine learning experts showcase how to use Amazon SageMaker. Learn more

Amazon SageMaker 10-Minute Deep Learning Model Tutorial

Feature Spotlight

Amazon SageMaker Ground Truth

Simplifying labeling workflows using Amazon SageMaker Ground Truth. Learn more

Accelerate Your Training Jobs Using Amazon FSx for Lustre

Speed up training on SageMaker with high-performance storage. Learn more

CloudShell

Feedback

© 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Training jobs | Amazon SageMaker

us-west-2.console.aws.amazon.com/sagemaker/home?region=us-west-2#/jobs?f0=kvm%253ACreationTimeAfter%253AFeb%2008%252C%202024%2003%253A50%20UTC

Services

Search

[Alt+S]

Oregon

voclabs/user3044095=cf0b5626-590a-11ee-a8a3-13a98f4cdb72 @ 08...

Amazon SageMaker

Getting started

Studio

Studio Lab

Canvas

RStudio

TensorBoard

Profiler

Admin configurations

Domains

Role manager

Images

Lifecycle configurations

SageMaker dashboard

Search

JumpStart

Foundation models

Computer vision models

Natural language processing models

Governance

HyperPod Clusters

Ground Truth

Amazon SageMaker > Training jobs

Training jobs

Info

Search training jobs

Creation time after : Feb 08, 2024 03:50 UTC

Actions

Create training job

< 1 >

	Name	Creation time	Duration	Job status	Warm pool status	Time left
	meta-textgeneration-llama-2-7b-2024-02-14-13-05-49-101	2/14/2024, 6:35:49 PM	15 minutes	Completed	-	-
	meta-textgeneration-llama-2-7b-2024-02-14-12-50-46-925	2/14/2024, 6:20:48 PM	14 minutes	Completed	-	-

CloudShell

Feedback

© 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences



meta-textgeneration-llama-2-7b

+

us-west-2.console.aws.amazon.com/sagemaker/home?region=us-west-2#/jobs/meta-textgeneration-llama-2-7b-2024-02-14-13-05-49-101

aws

Services

Search

[Alt+S]

Oregon

voclabs/user3044095=cf0b5626-590a-11ee-a8a3-13a98f4cdb72 @ 08...

Amazon SageMaker

Getting started

Studio

Studio Lab

Canvas

RStudio

TensorBoard

Profiler

Admin configurations

Domains

Role manager

Images

Lifecycle configurations

SageMaker dashboard

Search

JumpStart

Foundation models

Computer vision models

Natural language processing models

Governance

HyperPod Clusters

Ground Truth

Amazon SageMaker > Training jobs > meta-textgeneration-llama-2-7b-2024-02-14-13-05-49-101

meta-textgeneration-llama-2-7b-2024-02-14-13-05-49-101

CloneCreate model packageStopCreate model

Job settings

Job name	Status	SageMaker metrics time series	IAM role ARN
meta-textgeneration-llama-2-7b-2024-02-14-13-05-49-101	CompletedView history	Disabled	arn:aws:iam::087562234367:role/service-role/SageMaker-udacitySagemakerRole
ARN	Creation time	Training time (seconds)	
arn:aws:sagemaker:us-west-2:087562234367:training-job/meta-textgeneration-llama-2-7b-2024-02-14-13-05-49-101	Feb 14, 2024 13:05 UTC	778	
	Last modified time	Billable time (seconds)	
	Feb 14, 2024 13:20 UTC	778	
		Managed spot training savings	
		0%	
		Tuning job source/parent	
		-	

Algorithm

Algorithm ARN	Additional volume size (GB)	Maximum wait time for managed spot training(s)	Volume encryption key
-	30	-	-
Training image	Maximum runtime (s)	Managed spot training	
763104351884.dkr.ecr.us-west-2.amazonaws.com/huggingface-pytorch-training:2.0.0-transformers4.28.1-gpu-py310-cu118-ubuntu20.04	360000	Disabled	
Input mode			
File			

CloudShell

Feedback

© 2024, Amazon Web Services, Inc. or its affiliates. PrivacyTermsCookie preferences

meta-textgeneration-llama-2-7b

us-west-2.console.aws.amazon.com/sagemaker/home?region=us-west-2#/jobs/meta-textgeneration-llama-2-7b-2024-02-14-13-05-49-101

aws

Services

Search

[Alt+S]

Oregon

voclabs/user3044095=cf0b5626-590a-11ee-a8a3-13a98f4cdb72 @ 08...

Amazon SageMaker

Getting started

Studio

Studio Lab

Canvas

RStudio

TensorBoard

Profiler

Admin configurations

Domains

Role manager

Images

Lifecycle configurations

SageMaker dashboard

Search

JumpStart

Foundation models

Computer vision models

Natural language processing models

Governance

HyperPod Clusters

Ground Truth

Instance group	Instance type	Instance count	Keep alive period
-	ml.g5.2xlarge	1	-

Input data configuration: training

Channel name	Input mode	Data source	S3 data type
training	-	S3	S3Prefix
	Content type	Instance group	S3 data distribution type
	-	-	FullyReplicated
	Compression type		URI
	None		s3://genaiwithawsproject2024/training-datasets/finance
	Record wrapper type		
	None		

Input data configuration: code

Channel name	Input mode	Data source	S3 data type
code	-	S3	S3Prefix
	Content type	Instance group	S3 data distribution type
	-	-	FullyReplicated
	Compression type		URI
	None		s3://jumpstart-cache-prod-us-west-2/source-directory-tarballs/meta/transfer_learning/textgeneration/prepack/v1.0.3/sourcedir.tar.gz
	Record wrapper type		
	None		

Checkpoint configuration

CloudShell

Feedback

© 2024, Amazon Web Services, Inc. or its affiliates.

Privacy

Terms

Cookie preferences

meta-textgeneration-llama-2-7b

us-west-2.console.aws.amazon.com/sagemaker/home?region=us-west-2#/jobs/meta-textgeneration-llama-2-7b-2024-02-14-13-05-49-101

Services

Search

[Alt+S]

Oregon

voclabs/user3044095=cf0b5626-590a-11ee-a8a3-13a98f4cdb72 @ 08...

Amazon SageMaker

Getting started

Studio

Studio Lab

Canvas

RStudio

TensorBoard

Profiler

Admin configurations

Domains

Role manager

Images

Lifecycle configurations

SageMaker dashboard

Search

JumpStart

Foundation models

Computer vision models

Natural language processing models

Governance

HyperPod Clusters

Ground Truth

Name

huggingface-textgeneration:eval-loss

huggingface-textgeneration:eval-ppl

huggingface-textgeneration:train-loss

Regex

eval\_epoch\_loss=tensor\([0-9\.\+]

eval\_ppl=tensor\([0-9\.\+]

train\_epoch\_loss=[0-9\.\+]

Output data configuration

S3 output path

s3://sagemaker-us-west-2-087562234367/

Output encryption key

-

Hyperparameters

Key	Value
add_input_output_demarcation_key	"True"
chat_dataset	"False"
enable_fsdp	"True"
epoch	"5"
instruction_tuned	"False"
int8_quantization	"False"
learning_rate	"0.0001"
lora_alpha	"32"
lora_dropout	"0.05"
lora_r	"8"
max_input_length	"-1"
max_train_samples	"-1"
max_val_samples	"-1"
per_device_eval_batch_size	"1"
per_device_train_batch_size	"4"
preprocessing_num_workers	"None"
sagemaker_container_log_level	20
sagemaker_job_name	"meta-textgeneration-llama-2-7b-2024-02-14-13-05-49-101"
sagemaker_program	"transfer_learning.py"
sagemaker_region	"us-west-2"

CloudShell

Feedback

© 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences



meta-textgeneration-llama-2-7b

us-west-2.console.aws.amazon.com/sagemaker/home?region=us-west-2#/jobs/meta-textgeneration-llama-2-7b-2024-02-14-13-05-49-101

Services

Search

[Alt+S]

CloudShell

Feedback

Oregon

voclabs/user3044095=cf0b5626-590a-11ee-a8a3-13a98f4cdb72 @ 08...

Amazon SageMaker

Getting started

Studio

Studio Lab

Canvas

RStudio

TensorBoard

Profiler

Admin configurations

Domains

Role manager

Images

Lifecycle configurations

SageMaker dashboard

Search

JumpStart

Foundation models

Computer vision models

Natural language processing models

Governance

HyperPod Clusters

Ground Truth

S3 output path

s3://sagemaker-us-west-2-087562234367/

Output encryption key

-

Hyperparameters

Key	Value
add_input_output_demarcation_key	"True"
chat_dataset	"False"
enable_fsdp	"True"
epoch	"5"
instruction_tuned	"False"
int8_quantization	"False"
learning_rate	"0.0001"
lora_alpha	"32"
lora_dropout	"0.05"
lora_r	"8"
max_input_length	"-1"
max_train_samples	"-1"
max_val_samples	"-1"
per_device_eval_batch_size	"1"
per_device_train_batch_size	"4"
preprocessing_num_workers	"None"
sagemaker_container_log_level	20
sagemaker_job_name	"meta-textgeneration-llama-2-7b-2024-02-14-13-05-49-101"
sagemaker_program	"transfer_learning.py"
sagemaker_region	"us-west-2"
sagemaker_submit_directory	"/opt/ml/input/data/code/sourcedir.tar.gz"
seed	"10"
train_data_split_seed	"0"
validation_split_ratio	"0.2"

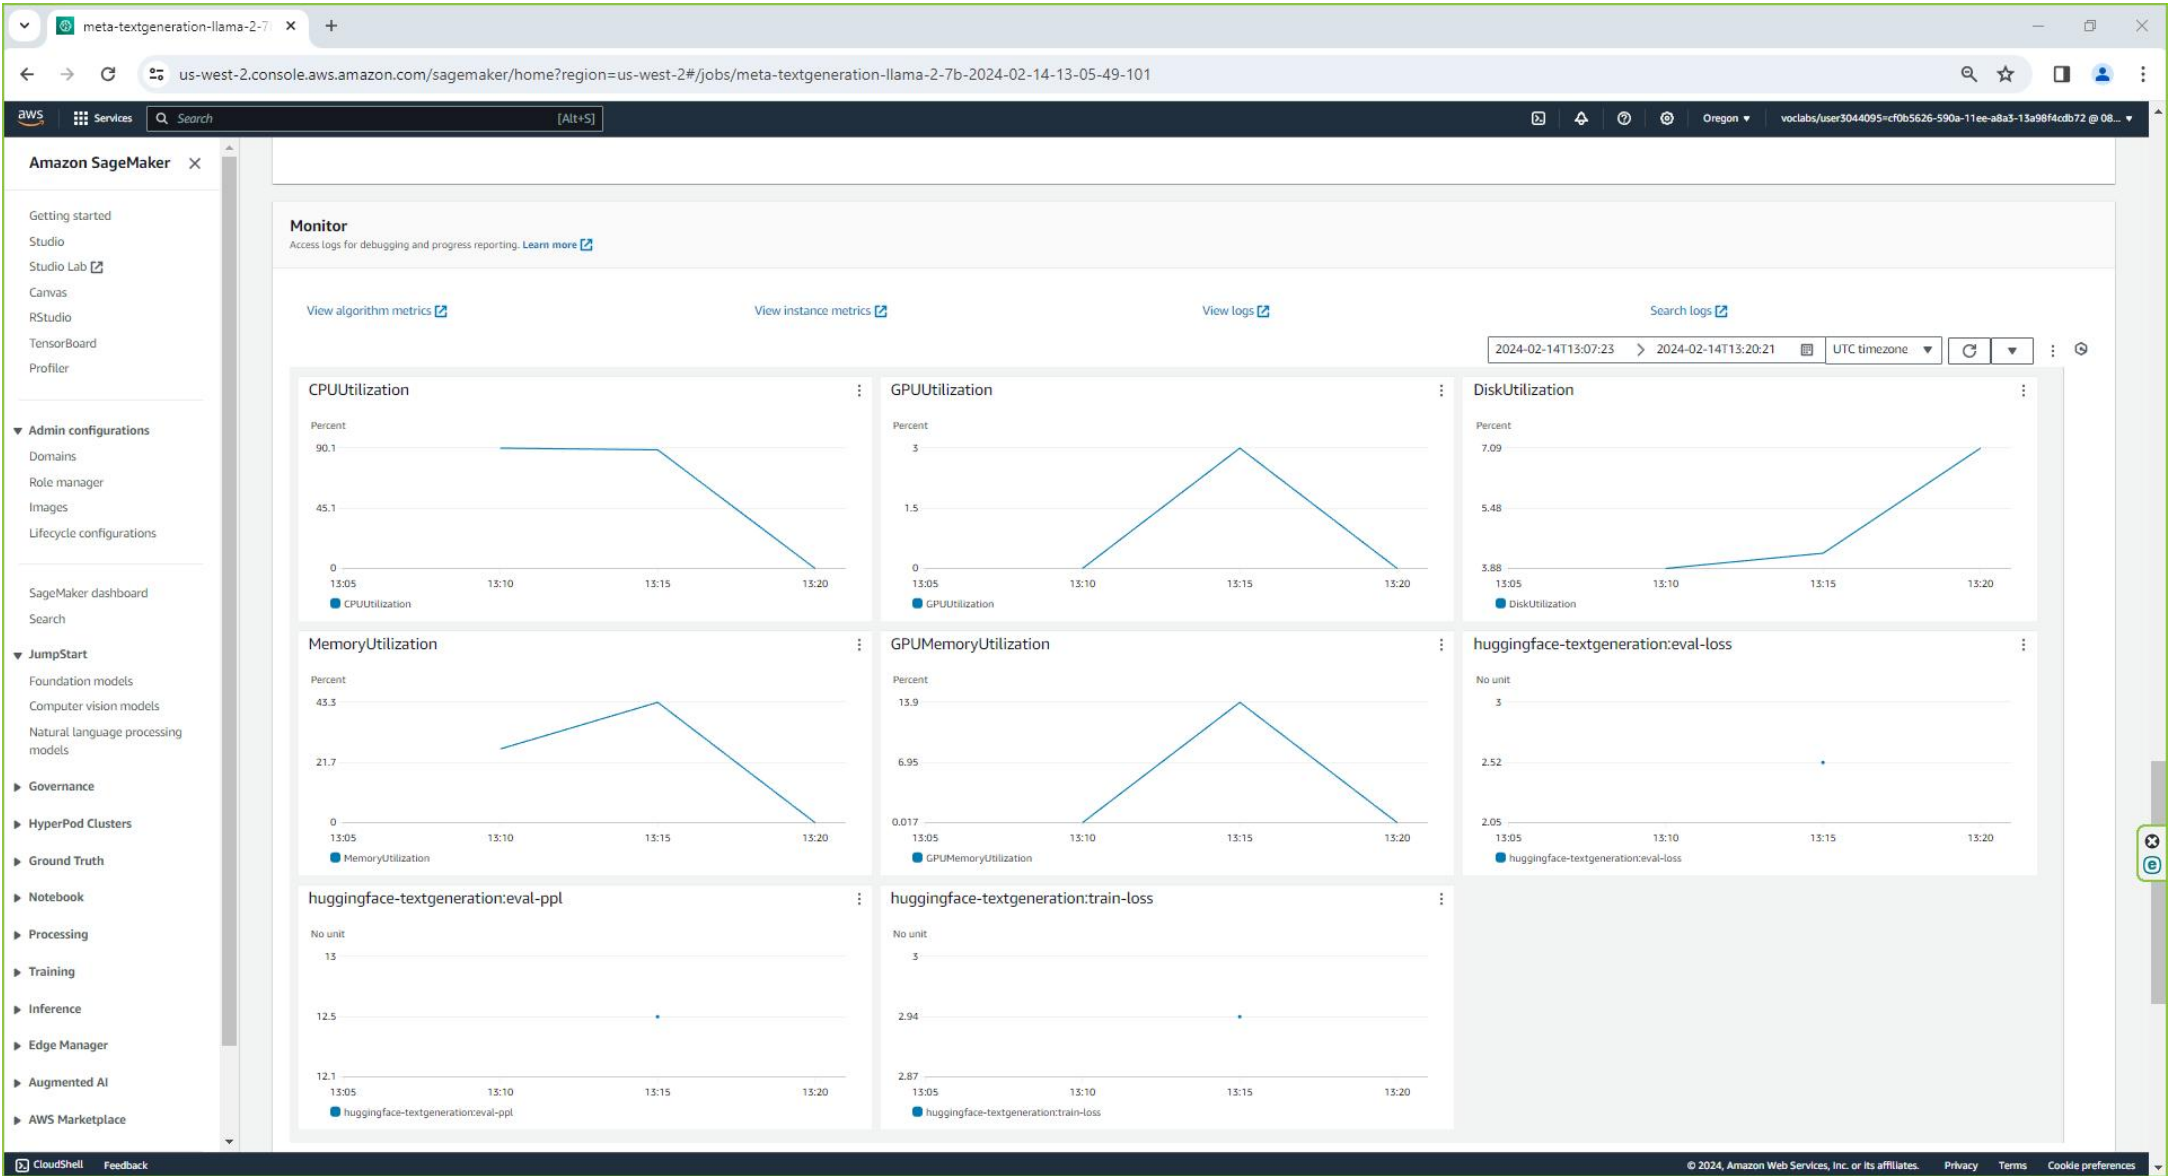
Network

© 2024, Amazon Web Services, Inc. or its affiliates.

Privacy

Terms

Cookie preferences





meta-textgeneration-llama-2-7b

us-west-2.console.aws.amazon.com/sagemaker/home?region=us-west-2#/jobs/meta-textgeneration-llama-2-7b-2024-02-14-13-05-49-101

ServicesSearch[Alt+S]

Amazon SageMaker

Getting started

Studio

Studio Lab

Canvas

RStudio

TensorBoard

Profiler

Admin configurations

Domains

Role manager

Images

Lifecycle configurations

SageMaker dashboard

Search

JumpStart

Foundation models

Computer vision models

Natural language processing models

Governance

HyperPod Clusters

Ground Truth

Output

S3 model artifact  
s3://sagemaker-us-west-2-087562234367/meta-textgeneration-llama-2-7b-2024-02-14-13-05-49-101/output/model

Environment variables

Key	Value
SageMakerGatedModelS3Uri	s3://jumpstart-private-cache-prod-us-west-2/meta-training/g5/v1.0.0/train-meta-textgeneration-llama-2-7b.tar.gz
accept_eula	true

Tags

Edit

Key	Value
sagemaker-sdk:jumpstart-model-id	meta-textgeneration-llama-2-7b
sagemaker-sdk:jumpstart-model-version	3.0.2

CloudShellFeedback

© 2024, Amazon Web Services, Inc. or its affiliates. PrivacyTermsCookie preferences

meta-textgeneration-llama-2-7b

+

us-west-2.console.aws.amazon.com/sagemaker/home?region=us-west-2#/jobs/meta-textgeneration-llama-2-7b-2024-02-14-12-50-46-925

aws

Services

Search

[Alt+S]

Oregon

voclabs/user3044095=cf0b5626-590a-11ee-a8a3-13a98f4cdb72 @ 08...

Amazon SageMaker

×

Getting started

Studio

Studio Lab

Canvas

RStudio

TensorBoard

Profiler

▼ Admin configurations

Domains

Role manager

Images

Lifecycle configurations

SageMaker dashboard

Search

▼ JumpStart

Foundation models

Computer vision models

Natural language processing models

► Governance

► HyperPod Clusters

► Ground Truth

► Notebook

► Processing

► Training

Amazon SageMaker

Training jobs

meta-textgeneration-llama-2-7b-2024-02-14-12-50-46-925

meta-textgeneration-llama-2-7b-2024-02-14-12-50-46-925

Clone

Create model package

Stop

Create model

Job settings

Job name

meta-textgeneration-llama-2-7b-2024-02-14-12-50-46-925

Status

Completed

View history

SageMaker metrics time series

Disabled

IAM role ARN

arn:aws:iam::087562234367:role/service-role/SageMaker-udacitySagemakerRole

ARN

arn:sagemaker:us-west-2:087562234367:training-job/meta-textgeneration-llama-2-7b-2024-02-14-12-50-46-925

Creation time

Feb 14, 2024 12:50 UTC

Training time (seconds)

773

Billable time (seconds)

773

Last modified time

Feb 14, 2024 13:05 UTC

Managed spot training savings

0%

Tuning job source/parent

-

Algorithm

Algorithm ARN

-

Additional volume size (GB)

30

Maximum wait time for managed spot training(s)

-

Volume encryption key

-

Training image

763104351884.dkr.ecr.us-west-2.amazonaws.com/huggingface-pytorch-training:2.0.0-transformers4.28.1-gpu-py310-cu118-ubuntu20.04

Maximum runtime (s)

360000

Managed spot training

Disabled

Input mode

File

Instance group

Instance type

Instance count

Keep alive period

-

ml.g5.2xlarge

1

-

CloudShell

Feedback

© 2024, Amazon Web Services, Inc. or its affiliates.

Privacy

Terms

Cookie preferences

meta-textgeneration-llama-2-7b

us-west-2.console.aws.amazon.com/sagemaker/home?region=us-west-2#/jobs/meta-textgeneration-llama-2-7b-2024-02-14-12-50-46-925

ServicesSearch[Alt+S]

CloudShellFeedback

© 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Amazon SageMaker

Getting started

Studio

Studio Lab

Canvas

RStudio

TensorBoard

Profiler

Admin configurations

Domains

Role manager

Images

Lifecycle configurations

SageMaker dashboard

Search

JumpStart

Foundation models

Computer vision models

Natural language processing models

Governance

HyperPod Clusters

Ground Truth

Input data configuration: training

Channel name	training	Input mode	-	Data source	S3	S3 data type	S3Prefix
		Content type	-	Instance group	-	S3 data distribution type	FullyReplicated
		Compression type	None			URI	s3://genaiwithawsproject2024/training-datasets/finance
		Record wrapper type	None				

Input data configuration: code

Channel name	code	Input mode	-	Data source	S3	S3 data type	S3Prefix
		Content type	-	Instance group	-	S3 data distribution type	FullyReplicated
		Compression type	None			URI	s3://jumpstart-cache-prod-us-west-2/source-directory-tarballs/meta/transfer_learning/textgeneration/prepare/v1.0.3/sourcedir.tar.gz
		Record wrapper type	None				

Checkpoint configuration

S3 output path	-	Local path	-
----------------	---	------------	---

meta-textgeneration-llama-2-7b

+

us-west-2.console.aws.amazon.com/sagemaker/home?region=us-west-2#/jobs/meta-textgeneration-llama-2-7b-2024-02-14-12-50-46-925

aws

Services

Search

[Alt+S]

🔍

🔔

?

🌐

Oregon

voclabs/user3044095-cf0b5626-590a-11ee-a8a3-13a98f4c6b72 @ 08...

Amazon SageMaker

×

Getting started

Studio

Studio Lab

Canvas

RStudio

TensorBoard

Profiler

▼ Admin configurations

Domains

Role manager

Images

Lifecycle configurations

▼ JumpStart

Foundation models

Computer vision models

Natural language processing models

► Governance

► HyperPod Clusters

► Ground Truth

► Notebook

► Processing

► Training

► Inference

► Edge Manager

► Augmented AI

S3 output path

-

Local path

-

Metrics

Name

huggingface-textgeneration:eval-loss

huggingface-textgeneration:eval-ppl

huggingface-textgeneration:train-loss

Regex

eval\_epoch\_loss=tensor\(\{[0-9\.]+\}

eval\_ppl=tensor\(\{[0-9\.]+\}

train\_epoch\_loss=\{[0-9\.]+\}

Output data configuration

S3 output path

s3://sagemaker-us-west-2-087562234367/

Output encryption key

-

Hyperparameters

Key

add\_input\_output\_demarcation\_key

chat\_dataset

enable\_fsdp

epoch

instruction\_tuned

int8\_quantization

learning\_rate

lora\_alpha

lora\_dropout

lora\_r

max\_input\_length

max\_train\_samples

max\_val\_samples

per\_device\_eval\_batch\_size

per\_device\_train\_batch\_size

preprocessing\_num\_workers

sagemaker\_container\_log\_level

sagemaker\_job\_name

sagemaker\_program

sagemaker\_region

sagemaker\_submit\_directory

seed

train\_data\_split\_seed

validation\_split\_ratio

Value

"True"

"False"

"True"

"5"

"False"

"False"

"0.0001"

"32"

"0.05"

"8"

"-1"

"-1"

"-1"

"1"

"4"

"None"

20

"meta-textgeneration-llama-2-7b-2024-02-14-12-50-46-925"

"transfer\_learning.py"

"us-west-2"

"/opt/ml/input/data/code/sourcedir.tar.gz"

"10"

"0"

"0.2"

CloudShell

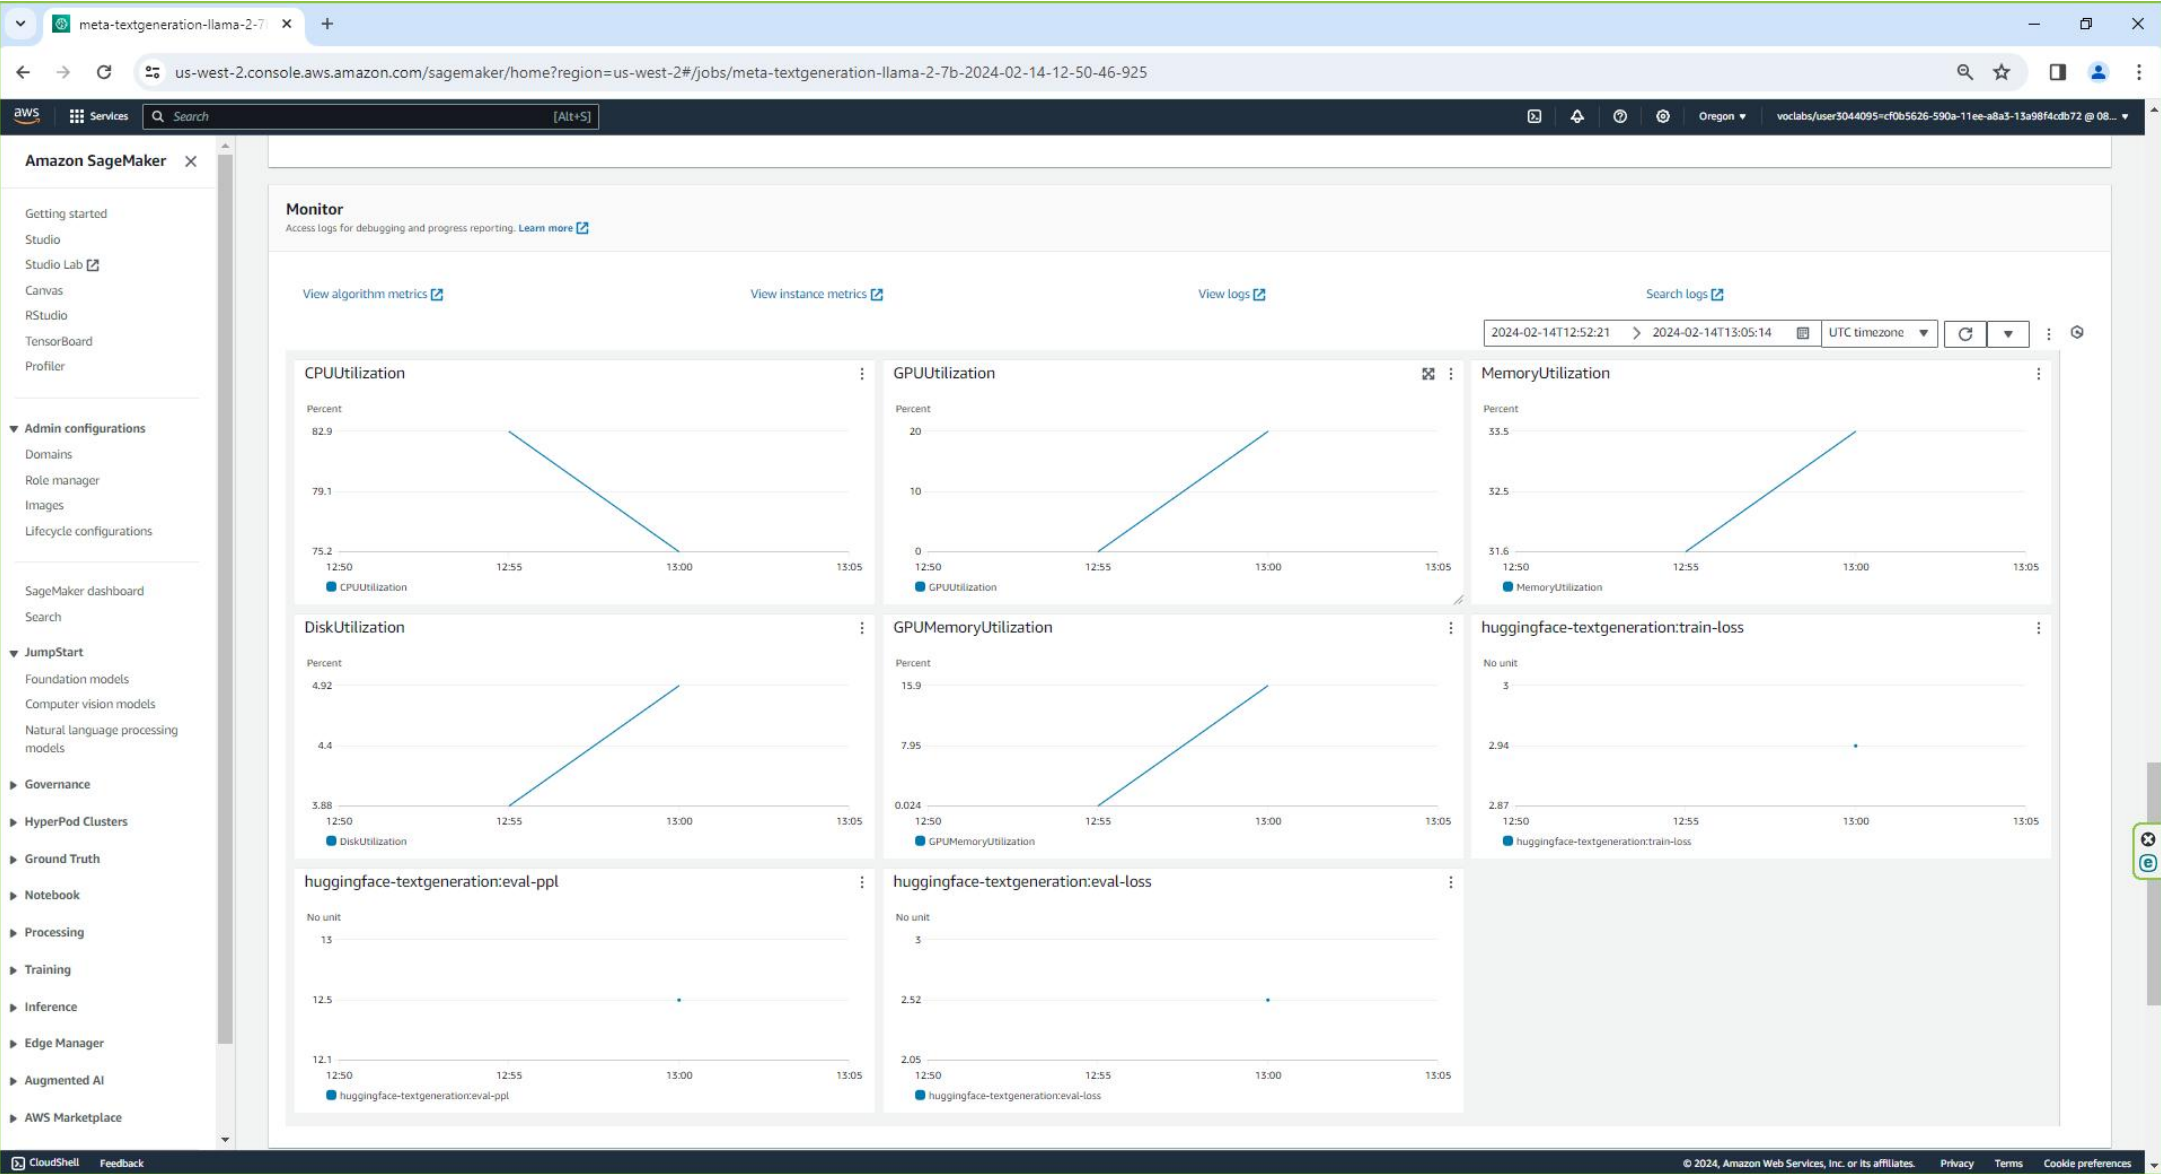
Feedback

© 2024, Amazon Web Services, Inc. or its affiliates.

Privacy

Terms

Cookie preferences



meta-textgeneration-llama-2-7b

us-west-2.console.aws.amazon.com/sagemaker/home?region=us-west-2#/jobs/meta-textgeneration-llama-2-7b-2024-02-14-12-50-46-925

aws

Services

Search

[Alt+S]

Oregon

voclabs/user3044095=cf0b5626-590a-11ee-a8a3-13a98f4cdb72 @ 08...

Amazon SageMaker

Getting started

Studio

Studio Lab

Canvas

RStudio

TensorBoard

Profiler

Admin configurations

Domains

Role manager

Images

Lifecycle configurations

SageMaker dashboard

Search

JumpStart

Foundation models

Computer vision models

Natural language processing models

Governance

HyperPod Clusters

Ground Truth

Output

S3 model artifact

s3://sagemaker-us-west-2-087562234367/meta-textgeneration-llama-2-7b-2024-02-14-12-50-46-925/output/model

Environment variables

Key	Value
SageMakerGatedModelS3Uri	s3://jumpstart-private-cache-prod-us-west-2/meta-training/g5/v1.0.0/train-meta-textgeneration-llama-2-7b.tar.gz
accept_eula	true

Tags

Edit

Key	Value
sagemaker-sdk:jumpstart-model-id	meta-textgeneration-llama-2-7b
sagemaker-sdk:jumpstart-model-version	3.0.2

CloudShell

Feedback

© 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences